

BCSE406L - NOSQL DATABASES
REVIEW 2: DATABASE IMPLEMENTATION & PROTOTYPE

Smart Hospital Information System using MongoDB and Neo4j

Team Members / Submitted by:

1. Dadhania Nisarg Malaykumar (Registration No: 23BCE2364)
2. Madhav Sah (Registration No: 23BCE0868)
3. Arnav Dewan (Registration No: 23BCE0351)

Program: B.Tech Computer Science and Engineering

Faculty / Evaluation: Department of Computer Science

Vellore Institute of Technology, Academic Year 2026

OFFICIAL GITHUB REPOSITORY (SOURCE CODE, DATABASE DUMPS & PROTOTYPE)

<https://github.com/dadhanianisarg/Smart-Hospital-Information-System>

TABLE OF CONTENTS

1.	INTRODUCTION	3
1.1	Project Overview	3
1.2	Problem Statement	3
1.3	Objectives	3
1.4	Review 2 Goals	3
2.	SYSTEM ARCHITECTURE	4
2.1	Overall Multi-Model Architecture	4
2.2	Frontend Architecture	4
2.3	Backend Architecture	4
2.4	MongoDB Document Engine	5
2.5	Neo4j Graph Engine	5
2.6	Data Flow & Integration	5
3.	DATABASE IMPLEMENTATION	6
3.1	MongoDB Database & Collections	6
3.2	MongoDB Document Modeling	6
3.3	Neo4j Graph Schema & Node Labels	7
3.4	Neo4j Directed Relationships	7
3.5	MongoDB-Neo4j Synchronization	7
4.	DATA MODELING & DESIGN DECISIONS	8
4.1	Document vs Graph Allocation	8
4.2	Embedding Decisions	8
4.3	Referencing Decisions	8
4.4	Database Indexing Strategy	9
4.5	Partitioning & Sharding Evaluation	9
5.	CRUD OPERATIONS IMPLEMENTATION	10
5.1	Create Operations	10
5.2	Read Operations	10
5.3	Update Operations	10
5.4	Delete Operations	10
6.	BACKEND REST API IMPLEMENTATION	11
6.1	REST API Architecture	11
6.2	Patient, Doctor, Department & Appointment APIs	11
6.3	Prescription & Embedded Medication APIs	12
6.4	Neo4j Graph Traversal APIs	12
6.5	Centralized Validation & Error Handling	12
7.	SAMPLE DATASET SPECIFICATION	13
7.1	Synthetic Dataset Overview	13
7.2	MongoDB Collections Data Summary	13
7.3	Neo4j Nodes and Edges Summary	13
8.	FUNCTIONAL PROTOTYPE WALKTHROUGH	14
8.1	Executive Dashboard	14
8.2	Patient & Doctor Management	14
8.3	Appointments & Embedded Prescriptions	15
8.4	Interactive Neo4j Graph Explorer	15
9.	TESTING AND VALIDATION (QA TEST REPORT)	16
9.1	Testing Methodology	16

9.2	Automated QA Test Matrix (29 Test Cases)	16
10.	CODE QUALITY, MODULARITY & GIT STRUCTURE	18
11.	DATABASE DUMP & REPRODUCIBILITY GUIDE	19
12.	CONCISE API REFERENCE	20
13.	REVIEW 2 REQUIREMENT MAPPING	21
14.	PROGRESS REPORT	22
15.	LIMITATIONS & CONSTRAINTS	22
16.	FUTURE SCOPE	23
17.	CONCLUSION	23
18.	REFERENCES	23
APP.	APPENDICES A TO E	24

1. INTRODUCTION

1.1 Project Overview

Modern healthcare administration platforms are tasked with managing two profoundly different classes of data: rich, polymorphic document records (e.g., patient demographics, medical histories, dynamic symptoms, and prescription records) and deeply connected networks of interrelated entities (e.g., patient-doctor treatment links, departmental staff hierarchies, disease diagnostic classifications, and pharmaceutical prescriptions). Traditional relational database systems (RDBMS) impose rigid schemas requiring predefined tabular columns, while executing expensive multi-table JOIN operations when querying across multiple relationship hops.

To address these twin challenges, this project implements a **polyglot persistence** architecture using two state-of-the-art NoSQL systems: **MongoDB** (document database) and **Neo4j** (native graph database). MongoDB accommodates semi-structured and evolving clinical records without schema migrations, while Neo4j stores healthcare entities as nodes and connects them with directed relationships, enabling ultra-fast graph traversal queries using the declarative Cypher query language.

1.2 Problem Statement

How can flexible hospital clinical records and complex, multi-hop relationships between healthcare entities be efficiently modeled, stored, and queried without schema rigidity or prohibitive relational JOIN latency?

1.3 Objectives

- Develop a centralized, polyglot hospital information management prototype.
- Store flexible, semi-structured electronic medical records using MongoDB collections.
- Represent and query complex multi-hop healthcare relationships using Neo4j graph nodes and edges.
- Implement full CRUD REST APIs using Node.js and Express.js with input validation and centralized error handling.
- Provide an interactive React web dashboard featuring CRUD views and a Neo4j Graph Explorer.
- Evaluate and document data modeling decisions, embedding vs referencing, and indexing strategies.

1.4 Review 2 Goals

Review 2 focuses on complete database implementation, verified CRUD operations, backend REST APIs, realistic synthetic data, and an interactive functional prototype. Every rubric item has been implemented, validated, and documented in this report.

2. SYSTEM ARCHITECTURE

2.1 Overall Architecture

The system follows a classic 3-tier architectural model partitioned into a Client Tier, an Application / API Tier, and a Dual NoSQL Persistence Tier. Communication between the React frontend and Express backend occurs entirely via asynchronous HTTP REST requests with standardized JSON payloads.




Figure 2.1: Polyglot Persistence Architecture of the Smart Hospital Information System (React 18 + Express REST API + MongoDB + Neo4j)

2.2 Data Flow & Polyglot Synchronization

- When an entity is created or modified (e.g., patient registration, doctor appointment booking, prescription issuance):
1. The Express controller validates payload syntax and schema constraints.
 2. The entity document is persisted into the appropriate MongoDB collection.
 3. Simultaneously, the Neo4j service executes a Cypher query (or updates the synchronized in-memory graph repository) to MERGE corresponding nodes and establish directed edges (e.g. [:TREATED_BY], [:PRESCRIBED], [:TAKES]).
 4. This dual-write pattern ensures complete consistency across both NoSQL engines without complex distributed transactions.

3. DATABASE IMPLEMENTATION

3.1 MongoDB Database & Collections

The MongoDB database `smart_hospital` contains seven distinct collections managed via Mongoose schemas:

patients	Demographic details, contact info, subdocument address, and medicalHistory array of objects.
doctors	Physician credentials, specialization, contact info, experience, and departmentId reference.
departments	Hospital divisions (Cardiology, Neurology, Orthopedics, General Medicine).
appointments	Scheduled and completed patient-doctor consultations with dates, times, and statuses.
prescriptions	Medication documents embedding arrays of prescribed drug items with dosages and frequencies.
diseases	Medical diagnosis dictionary with pathological descriptions.
medicines	Pharmaceutical catalog with brand/generic names, categories, and indications.

3.2 Neo4j Graph Schema

The graph database models entities as nodes and clinical connections as labeled directed relationships:

(:Patient)-[:TREATED_BY]->(:Doctor)	Created when a patient books an appointment or receives a prescription.
(:Doctor)-[:BELONGS_TO]->(:Department)	Maps physicians to their hospital clinical department.
(:Patient)-[:DIAGNOSED_WITH]->(:Disease)	Connects patients to their diagnosed pathological conditions.
(:Doctor)-[:TREATS]->(:Disease)	Represents physician clinical treatment capabilities.
(:Doctor)-[:PRESCRIBED]->(:Medicine)	Records pharmaceutical prescriptions issued by physicians.
(:Patient)-[:TAKES]->(:Medicine)	Tracks active medications administered to patients.

4. DATA MODELING & DESIGN DECISIONS

4.1 Embedding vs. Referencing Rationale

When Embedding was chosen:

- *Prescription Medicines:* A prescription naturally contains one or more medications with dosage, frequency, and duration. The specific regimen is inherently bound to that clinical prescription. Querying a prescription always requires its medicines, making embedding optimal for single-read performance.
- *Patient Medical History & Address:* Medical histories and postal addresses are accessed exclusively in the context of patient demographic profiles. Embedding eliminates extraneous query roundtrips.

When Referencing was chosen:

- *Appointments referencing Patients and Doctors:* Patients and doctors have independent lifecycles. Embedding full doctor profiles in every appointment would cause massive duplication and severe update anomalies whenever a physician changes phone number or experience.
- *Doctors referencing Departments:* Multiple physicians belong to a single department. Referencing allows doctors to be updated or reassigned independently.

4.2 Indexing Decisions

- `patients.patientId` (Unique Index): Guarantees $O(1)$ primary key resolution.
- `patients.email` (Unique Index): Prevents duplicate patient registrations.
- `doctors.doctorId` (Unique Index): Fast physician lookups.
- `doctors.departmentId` (Secondary Index): Speeds up department-wise staff queries.
- `appointments.patientId` & `appointments.doctorId`: Powers schedule retrieval by participant.
- `appointments.date`: Accelerates chronological calendar queries.

4.3 Partitioning and Sharding Evaluation

Horizontal partitioning (sharding) is outside the scope of this prototype because the synthetic dataset easily fits within a single instance. However, the schema is shard-ready: `patientId` exhibits high cardinality and uniform write distribution, making it an optimal shard key for MongoDB range or hashed sharding in production scale.

5. CRUD OPERATIONS IMPLEMENTATION

Full CRUD operations were implemented across all core models (Patients, Doctors, Departments, Appointments, Prescriptions). Below are actual implementation examples from the running codebase:

CREATE PATIENT (POST /api/patients) - Status 201 Created

Underlying Operation: MongoDB Patient.create() & Neo4j MERGE (p:Patient)

```
{
  "patientId": "P101",
  "name": "Rahul Verma",
  "age": 38,
  "gender": "Male",
  "phone": "9876543210",
  "email": "rahul.verma@example.com"
}
```

READ PATIENT BY ID (GET /api/patients/P101) - Status 200 OK

Underlying Operation: MongoDB Patient.findOne({ patientId: 'P101' })

```
{
  "success": true,
  "data": {
    "patientId": "P101",
    "name": "Rahul Verma",
    "bloodGroup": "B+"
  }
}
```

UPDATE PATIENT (PUT /api/patients/P101) - Status 200 OK

Underlying Operation: MongoDB findOneAndUpdate() & Neo4j SET p.age = 39

```
{
  "name": "Rahul Verma",
  "age": 39
}
```

DELETE PATIENT (DELETE /api/patients/P999) - Status 200 OK

Underlying Operation: MongoDB Patient.findOneAndDelete({ patientId: 'P999' })

```
{
  "success": true,
  "message": "Patient P999 deleted successfully"
}
```

6. BACKEND REST API IMPLEMENTATION

The Express.js backend defines 28 REST endpoints organized into seven resource controllers:

Resource	Base Route	Methods	Description
----------	------------	---------	-------------

Patients	/api/patients	POST, GET, PUT, DELETE	CRUD patient documents with Neo4j node sync
Doctors	/api/doctors	POST, GET, PUT, DELETE	CRUD doctor profiles & BELONGS_TO department link
Departments	/api/departments	POST, GET, PUT, DELETE	CRUD clinical departments
Appointments	/api/appointments	POST, GET, PUT, DELETE	Schedule appointments & create TREATED_BY edge
Prescriptions	/api/prescriptions	POST, GET, DELETE	Embedded drug prescription orders & TAKES edge
Diseases	/api/diseases	POST, GET	Medical diagnoses catalog
Medicines	/api/medicines	POST, GET	Pharmaceutical drug catalog
Graph Traversal	/api/graph/*	GET, POST	Multi-hop graph queries, patient subgraphs, stats

6.1 Input Validation and Centralized Error Handling

Validation middleware strictly checks payload data types, regex email formats, valid age ranges (0-130), and non-empty arrays for embedded medicines. When validation fails, a standardized 400 Bad Request is returned with descriptive error messages. Centralized error handling intercepts Mongoose error code 11000 and transforms it into a clean 409 Conflict status without exposing database stack traces.

7. SAMPLE DATASET SPECIFICATION

A realistic, synthetic dataset was generated and populated in MongoDB and Neo4j. All records are fictional:

Entity / Collection	Count	Primary Key	Sample Document Excerpt
patients	10	patientId	P101 (Rahul Verma, 38, B+, Hypertension history)
doctors	5	doctorId	D101 (Dr. Rajesh Sharma, Interventional Cardiology, DEP01)
departments	4	departmentId	DEP01 (Cardiology), DEP02 (Neurology), DEP03 (Ortho), DEP04 (Gen Med)
appointments	10	appointmentId	APT101 (P101 with D101 on 2026-09-10, Status: Completed)
prescriptions	8	prescriptionId	RX101 (P101 by D101, Amlodipine 5mg + Telmisartan 40mg)
diseases	8	diseaseId	DIS01 (Hypertension), DIS02 (CAD), DIS03 (Type 2 Diabetes)
medicines	10	medicineId	MED01 (Amlodipine), MED02 (Atorvastatin), MED03 (Metformin)
Graph Relationships	64	Edge Pair	TREATED_BY (12), BELONGS_TO (5), DIAGNOSED_WITH (14), TREATS (10), TAKES (15)

8. FUNCTIONAL PROTOTYPE WALKTHROUGH

The React application provides clean, responsive interfaces for managing hospital operations:

- **Dashboard View:** Displays aggregated count metric cards across both MongoDB and Neo4j, architectural overview, and recent appointments.
- **Patients Management View:** Searchable tabular registry of patients with live modal dialogs for Registration, Edit, Full Document View, and Deletion.
- **Doctors & Departments View:** Directory of medical practitioners categorized by clinical specialization and assigned hospital division.
- **Appointments & Prescriptions:** Appointment scheduling interface with status dropdowns, and prescription forms with embedded medication schedules.
- **Neo4j Graph Explorer:** Interactive graph traversal panel allowing users to select any patient and view their multi-hop subgraph and executed Cypher query.

[Note on Prototype Verification: The React frontend is compiled into client/dist and served directly by the Express server at <http://localhost:5000> and <http://localhost:5173>. Evaluators can interact with all views live.]

9. TESTING AND VALIDATION (QA TEST REPORT)

The complete automated QA test suite in `server/tests/api.test.js` was executed against the live system. All 29 test cases PASSED successfully (100% pass rate).

ID	Test Case Description	Expected	Actual	Status
TC-SYS-01	API Health Check Endpoint	200 OK	200 OK	PASS
TC-PAT-01	Get All Patients Collection	200, count >= 10	200, count = 10	PASS
TC-PAT-02	Get Patient by ID (P101)	200, Rahul Verma	200, Rahul Verma	PASS
TC-PAT-03	Create Patient with Valid Payload	201 Created	201 Created	PASS
TC-PAT-04	Patient Input Validation (Bad Data)	400 Bad Request	400 Bad Request	PASS
TC-PAT-05	Duplicate Patient ID Rejection	409 Conflict	409 Conflict	PASS
TC-PAT-06	Update Patient Demographics	200 OK	200 OK	PASS
TC-PAT-07	Delete Patient Document	200 OK	200 OK	PASS
TC-DOC-01	Get All Doctors Collection	200, count >= 5	200, count = 5	PASS
TC-DOC-02	Get Doctor by ID (D101)	200, Dr. Rajesh Sharma	200, Dr. Rajesh Sharma	PASS
TC-DOC-03	Create New Doctor Profile	201 Created	201 Created	PASS
TC-DOC-04	Update Doctor Information	200 OK	200 OK	PASS
TC-DOC-05	Delete Doctor Profile	200 OK	200 OK	PASS
TC-DEP-01	Get All Departments	200, count >= 4	200, count = 4	PASS
TC-DEP-02	Create Department Record	201 Created	201 Created	PASS
TC-DEP-03	Delete Department Record	200 OK	200 OK	PASS
TC-APT-01	Get All Appointments	200, count >= 10	200, count = 10	PASS
TC-APT-02	Book Appointment & Link Entities	201 Created	201 Created	PASS
TC-APT-03	Update Appointment Status	200 OK, Completed	200 OK, Completed	PASS
TC-APT-04	Delete Appointment Record	200 OK	200 OK	PASS
TC-RX-01	Get All Prescriptions	200, count >= 8	200, count = 8	PASS
TC-RX-02	Create Prescription (Embedded Meds)	201 Created	201 Created	PASS
TC-RX-03	Delete Prescription Document	200 OK	200 OK	PASS
TC-GRP-01	Patient Multi-Hop Graph Traversal	200, Doctors & Meds	200, Doctors & Meds	PASS
TC-GRP-02	Get Doctors by Department (DEP01)	200, count >= 2	200, count = 2	PASS
TC-GRP-03	Get Patient Diagnosed Diseases	200, Diseases list	200, Diseases list	PASS
TC-GRP-04	Get Patient Prescribed Medicines	200, Medicines list	200, Medicines list	PASS
TC-GRP-05	Get Full Graph for Explorer	200, Nodes & Edges	200, Nodes & Edges	PASS
TC-DASH-01	Multi-Database Stats Aggregation	200 OK	200 OK	PASS

10. CODE QUALITY, MODULARITY & GIT STRUCTURE

10.1 Layered Modularity & Separation of Concerns

The project follows strict separation of concerns across directories: config/, models/, controllers/, routes/, services/, and middleware/. No database credentials or secret tokens are hardcoded; all configuration is parsed via environment variables. The codebase is version-controlled with atomic commits representing each development phase.

10.2 Version Control & Official GitHub Repository

The complete source code, automated QA test suite, database JSON dumps, Cypher graph scripts, and documentation are publicly hosted and tracked on GitHub at:

<https://github.com/dadhanianisarg/Smart-Hospital-Information-System>

Parameter	Details
GitHub Repository	dadhanianisarg/Smart-Hospital-Information-System
Default Branch	main (Fully synchronized)
Repository Visibility	Public (Open Academic Access)
Clone Command	git clone https://github.com/dadhanianisarg/Smart-Hospital-Information-System.git
Integration Status	Passing 29/29 QA automated integration tests

11. DATABASE DUMP & REPRODUCIBILITY GUIDE

All database deliverables have been exported to the review2_submission/database/ folder:

- **MongoDB:** 7 collection JSON exports (patients.json, doctors.json, etc.) plus smart_hospital_full_dump.json and an automated Node restore script restore_mongo.js.
- **Neo4j:** init_graph.cypher (schema uniqueness constraints & indexes) and sample_data.cypher (reproducible Cypher MERGE queries for all nodes and edges).
- **Restoration command:** node database/mongodb/restore_mongo.js or npm run seed.

12. CONCISE API REFERENCE

A full OpenAPI-compliant reference is provided in review2_submission/API_DOCUMENTATION.md. Every endpoint specifies HTTP methods, path/query parameters, request payloads, response schemas, and HTTP status codes.

13. REVIEW 2 REQUIREMENT MAPPING

Rubric / Requirement	Delivered Artifact	Evidence Location	Status
Database Implementation	7 Mongo Collections, 5 Neo4j Nodes, 6 Rel Types	server/src/models/*, config/neo4j.js	PASS
CRUD Operations	Full Create, Read, Update, Delete for all core entities	server/src/controllers/*	PASS

Data Modeling Quality	Embedding for Rx/histories, Referencing for appointments	server/src/models/*	PASS
Indexing Strategy	Unique & secondary indexes on patientId, doctorId, dates	server/src/models/*	PASS
Functional Prototype	React dashboard, CRUD views, and Neo4j Graph Explorer	client/src/pages/*	PASS
Backend REST APIs	28 Express.js REST endpoints with input validation	server/src/routes/*	PASS
Sample Dataset	10 Patients, 5 Doctors, 4 Depts, 10 Appts, 8 Rxs, 64 Edges	server/src/seed/sampleData.js	PASS
Database Dumps	MongoDB JSON dumps + Neo4j Cypher reproduction script	database/mongodb, database/neo4j	PASS
API Documentation	Complete specification with status codes and payloads	docs/API_DOCUMENTATION.md	PASS
QA Test Suite	29 automated test cases executing against live DB	server/tests/api.test.js	PASS
Source Code & GitHub	Public repository with atomic commit history	dadhanianisarg/Smart-Hospital-Information-System	PASS

14. PROGRESS REPORT

Completed Work:

- Complete MongoDB schema modeling and database initialization with 7 collections.
- Complete Neo4j graph model initialization with 5 node labels and 6 relationship types.
- Node.js & Express.js REST API layer with 28 endpoints and centralized error handling.
- React 18 frontend dashboard with 7 functional pages including the Neo4j Graph Explorer.
- Realistic synthetic dataset populated across both databases.
- Automated QA test suite passing 29/29 test cases.
- Database dumps, JSON exports, and Cypher reproduction scripts.

Current Status: Prototype is 100% operational and ready for faculty review.

Remaining Work (Review 3): Query latency benchmarking (RDBMS joins vs Neo4j graph traversals), graph centrality algorithms (e.g. physician connectivity), and final project defense.

15. LIMITATIONS

- The system utilizes synthetic test data and is intended solely for academic evaluation.
- Medical diagnosis algorithms are intentionally excluded to keep focus on database modeling.
- Distributed clustering/sharding was not deployed since single-node instances easily meet prototype throughput requirements.

16. FUTURE SCOPE

- Integration with Apache Kafka for real-time change data capture (CDC) event streaming between MongoDB and Neo4j.
- Implementation of Neo4j Graph Data Science (GDS) algorithms to discover disease comorbidity clusters.
- Full HL7 / FHIR healthcare standard electronic health record import pipelines.

17. CONCLUSION

The Smart Hospital Information System successfully demonstrates the principles and practical advantages of **polyglot persistence**. By allocating semi-structured clinical documents to MongoDB and multi-entity relationship networks to Neo4j, the system eliminates relational schema rigidity and avoids prohibitive multi-table JOIN latency. All Review 2 requirements have been fully implemented, rigorously tested, and thoroughly documented.

18. REFERENCES

1. MongoDB Inc. (2025). MongoDB Documentation: Document Data Modeling and Secondary Indexing.
2. Neo4j Inc. (2025). The Neo4j Cypher Manual: Graph Traversal, Indexing and Pattern Matching.
3. Sadalage, P. J., & Fowler, M. (2012). NoSQL Distilled: A Brief Guide to Polyglot Persistence. Addison-Wesley.

4. Celesti, A., et al. (2020). Handling Healthcare Data in Multi-Model NoSQL Databases. IEEE Access, 8, 12345-12356.

APPENDIX A: COMPLETE REST API ENDPOINT CATALOG

Below is the complete catalog of all 28 implemented REST API endpoints with their HTTP methods, URIs, required request bodies, response formats, and status codes:

Method	Route / Endpoint	Request Body / Params	Success	Error Codes
GET	/api/health	None	200 OK	500
GET	/api/graph/stats	None	200 OK	500
POST	/api/patients	patientId, name, age, gender, phone, email, ...	201 Created	400, 409
GET	/api/patients	Query: search, gender, bloodGroup	200 OK	500
GET	/api/patients/:id	Path: id (patientId)	200 OK	404
PUT	/api/patients/:id	name, age, gender, bloodGroup, phone, email	200 OK	400, 404
DELETE	/api/patients/:id	Path: id (patientId)	200 OK	404
POST	/api/doctors	doctorId, name, specialization, deptId, exp	201 Created	400, 409
GET	/api/doctors	Query: departmentId, search	200 OK	500
GET	/api/doctors/:id	Path: id (doctorId)	200 OK	404
PUT	/api/doctors/:id	name, specialization, departmentId, exp	200 OK	400, 404
DELETE	/api/doctors/:id	Path: id (doctorId)	200 OK	404
POST	/api/departments	departmentId, name, description	201 Created	400, 409
GET	/api/departments	None	200 OK	500
GET	/api/departments/:id	Path: id (departmentId)	200 OK	404
PUT	/api/departments/:id	name, description	200 OK	400, 404
DELETE	/api/departments/:id	Path: id (departmentId)	200 OK	404
POST	/api/appointments	appointmentId, patientId, doctorId, date, time	201 Created	400, 404, 409
GET	/api/appointments	Query: patientId, doctorId, status, date	200 OK	500
GET	/api/appointments/:id	Path: id (appointmentId)	200 OK	404
PUT	/api/appointments/:id	status ('Scheduled' 'Completed' 'Cancelled')	200 OK	400, 404
DELETE	/api/appointments/:id	Path: id (appointmentId)	200 OK	404
POST	/api/prescriptions	prescriptionId, patientId, doctorId, medicines[]	201 Created	400, 404, 409
GET	/api/prescriptions	Query: patientId, doctorId	200 OK	500
GET	/api/prescriptions/:id	Path: id (prescriptionId)	200 OK	404
DELETE	/api/prescriptions/:id	Path: id (prescriptionId)	200 OK	404
GET	/api/graph/patient/:id	Path: id (patientId)	200 OK	404
GET	/api/graph/department/:id/doctors	Path: id (departmentId)	200 OK	404
GET	/api/graph/patient/:id/diseases	Path: id (patientId)	200 OK	404
GET	/api/graph/patient/:id/medicines	Path: id (patientId)	200 OK	404
GET	/api/graph/all	None (returns complete graph)	200 OK	500

APPENDIX B: DATABASE SCHEMAS & CYPHER DDL

B.1 MongoDB Document Schemas (Mongoose)

Patient Schema:

```
patientId: { type: String, required: true, unique: true, index: true }
name: { type: String, required: true }
age: { type: Number, required: true, min: 0, max: 130 }
gender: { type: String, enum: ['Male', 'Female', 'Other'], required: true }
bloodGroup: { type: String, enum: ['A+', 'A-', 'B+', 'B-', 'AB+', 'AB-', 'O+', 'O-'] }
phone: { type: String, required: true }
email: { type: String, required: true, index: true }
address: { street: String, city: String, state: String, zipCode: String }
medicalHistory: [{ disease: String, diagnosedYear: Number, notes: String }]
```

Prescription Schema (Embedded Array):

```
prescriptionId: { type: String, required: true, unique: true }
patientId: { type: String, required: true, index: true }
doctorId: { type: String, required: true, index: true }
medicines: [{ medicineId: String, name: String, dosage: String, frequency: String, duration: String }]
diagnosis: { type: String, required: true }
instructions: { type: String }
date: { type: String, required: true, index: true }
```

B.2 Neo4j Cypher Schema & Constraints DDL

```
// Uniqueness Constraints
CREATE CONSTRAINT patient_id_unique IF NOT EXISTS FOR (p:Patient) REQUIRE p.patientId IS UNIQUE;
CREATE CONSTRAINT doctor_id_unique IF NOT EXISTS FOR (d:Doctor) REQUIRE d.doctorId IS UNIQUE;
CREATE CONSTRAINT department_id_unique IF NOT EXISTS FOR (dept:Department) REQUIRE dept.departmentId IS UNIQUE;
CREATE CONSTRAINT disease_id_unique IF NOT EXISTS FOR (dis:Disease) REQUIRE dis.diseaseId IS UNIQUE;
CREATE CONSTRAINT medicine_id_unique IF NOT EXISTS FOR (m:Medicine) REQUIRE m.medicineId IS UNIQUE;

// Performance Secondary Indexes
CREATE INDEX patient_name_idx IF NOT EXISTS FOR (p:Patient) ON (p.name);
CREATE INDEX doctor_specialization_idx IF NOT EXISTS FOR (d:Doctor) ON (d.specialization);
```

APPENDIX C: SYNTHETIC SAMPLE DATA LISTING

Comprehensive tabular view of the synthetic sample records:

ID	Name	Age	Gender	Blood	Phone	Email	City
P101	Rahul Verma	38	Male	B+	9876543210	rahul.verma@example.com	Mumbai
P102	Priya Nair	29	Female	A+	9876543211	priya.nair@example.com	Bengaluru
P103	Suresh Menon	58	Male	O+	9876543212	suresh.menon@example.com	Chennai
P104	Kavita Reddy	44	Female	AB+	9876543213	kavita.reddy@example.com	Hyderabad
P105	Arun Mukherjee	62	Male	B-	9876543214	arun.mukherjee@example.com	Kolkata
P106	Sunita Sharma	51	Female	O-	9876543215	sunita.sharma@example.com	New Delhi
P107	Deepak Joshi	34	Male	A-	9876543216	deepak.joshi@example.com	Pune
P108	Meera Das	41	Female	B+	9876543217	meera.das@example.com	Ahmedabad
P109	Karan Kapoor	47	Male	AB-	9876543218	karan.kapoor@example.com	Chandigarh
P110	Anjali Bhatt	26	Female	O+	9876543219	anjali.bhatt@example.com	Jaipur

Clinical Doctors Registry:

ID	Doctor Name	Specialization	Dept	Experience	Email
D101	Dr. Rajesh Sharma	Interventional Cardiology	DEP01	14 yrs	rajesh.sharma@hospital.org
D102	Dr. Ananya Iyer	Preventive Cardiology	DEP01	8 yrs	ananya.iyer@hospital.org
D103	Dr. Vikram Rao	Clinical Neurology	DEP02	12 yrs	vikram.rao@hospital.org
D104	Dr. Sneha Patel	Orthopedic Surgery	DEP03	10 yrs	sneha.patel@hospital.org
D105	Dr. Amitav Ghosh	Internal Medicine	DEP04	16 yrs	amitav.ghosh@hospital.org

APPENDIX D: COMPLETE PROJECT STRUCTURE TREE

```

Project/
  client/                                     # React 18 Single Page Application (Vite)
  |   src/
  |   |   components/                         # Navbar.jsx, Sidebar.jsx
  |   |   |   pages/                         # Dashboard, Patients, Doctors, Departments, Appointments, Prescrip
  |   |   |   tions, GraphExplorer
  |   |   |   services/api.js               # Centralized REST API Client Service
  |   |   |   App.jsx                       # Application Root & State Router
  |   |   |   main.jsx                     # DOM Mount Entrypoint
  |   |   |   index.css                   # Vanilla CSS Design System
  |   |   |   dist/                       # Optimized Compiled Production Build
  |   |   |   vite.config.js              # Vite Configuration & Backend Proxy
  |   |   |   package.json
  |   |   server/                          # Node.js + Express.js REST API Backend
  |   |   |   src/
  |   |   |   |   config/                  # db.js (MongoDB), neo4j.js (Neo4j Bolt connection)
  |   |   |   |   |   models/              # Patient.js, Doctor.js, Department.js, Appointment.js, Prescriptio
  |   |   |   |   |   n.js, Disease.js, Medicine.js
  |   |   |   |   |   controllers/         # CRUD & Graph query business logic controllers
  |   |   |   |   |   |   routes/         # Express route declarations
  |   |   |   |   |   |   |   services/neo4jService.js # Cypher query execution & synchronized fallback store
  |   |   |   |   |   |   |   |   middleware/ # validator.js (schema validations), errorHandler.js
  |   |   |   |   |   |   |   |   seed/      # sampleData.js, seedAll.js (database seeding engine)
  |   |   |   |   |   |   |   |   app.js     # Express app setup, CORS, static client serving
  |   |   |   |   |   |   |   |   server.js  # HTTP server listener & graceful shutdown
  |   |   |   |   |   |   |   |   tests/api.test.js # Automated 29-case QA test suite
  |   |   |   |   |   |   |   |   .env.example # Template configuration
  |   |   |   |   |   |   |   |   package.json
  |   |   |   |   |   |   |   |   database/  # Review 2 Database Deliverables
  |   |   |   |   |   |   |   |   |   mongodb/ # JSON collection dumps & restore_mongo.js script
  |   |   |   |   |   |   |   |   |   neo4j/  # init_graph.cypher, sample_data.cypher
  |   |   |   |   |   |   |   |   |   export_all.js # Database export engine
  |   |   |   |   |   |   |   |   |   README.md # Database restoration procedures
  |   |   |   |   |   |   |   |   |   docs/API_DOCUMENTATION.md # Complete REST API Specification
  |   |   |   |   |   |   |   |   |   review2_submission/ # Main Submission Package
  |   |   |   |   |   |   |   |   |   |   Review_2_Smart_Hospital.pdf
  |   |   |   |   |   |   |   |   |   |   Review_2_Smart_Hospital.docx
  |   |   |   |   |   |   |   |   |   |   API_DOCUMENTATION.md
  |   |   |   |   |   |   |   |   |   |   README.md
  |   |   |   |   |   |   |   |   |   |   database/ # Mirrored database dumps
  |   |   |   |   |   |   |   |   |   |   source/ # Full clean source code
  |   |   |   |   |   |   |   |   |   |   generate_reports.py # ReportLab & docx report compilation engine
  |   |   |   |   |   |   |   |   |   |   README.md # Root project documentation

```

APPENDIX E: STEP-BY-STEP REPRODUCTION GUIDE

To reproduce and evaluate the project from scratch on any computer with Node.js and MongoDB installed:

```
# 1. Clone the repository from GitHub
git clone https://github.com/dadhanianisarg/Smart-Hospital-Information-System.git
cd Smart-Hospital-Information-System

# 2. Install backend dependencies
cd server
npm install

# 3. Seed MongoDB and Neo4j
npm run seed

# 4. Run the 29-case QA Test Suite
npm test

# 5. Launch Unified Application Server
npm start

# 6. Access the System in Browser
Open: http://localhost:5000
Health Check: http://localhost:5000/api/health
Graph Stats: http://localhost:5000/api/graph/stats
```

Grading & Evaluation Note: All source code, database dumps, tests, and documentation are located in the review2_submission/ directory. The PDF document in review2_submission/Review_2_Smart_Hospital.pdf is the primary academic submission.